

MSiA-413 Introduction to Databases and Information Retrieval

Lecture 10 Indexes

Instructor: Nikos Hardavellas

Slides adapted from Steve Tarzia, Silberschatz

Last Lecture: introduced the *storage hierarchy*

Larger, but slower ↓

delay		capacity
0.3ns	CPU Registers	1 kB (kilobyte)
5ns	CPU Caches (L2)	16 MB
50ns	Random Access Memory (RAM)	16 GB
100 μ s	Flash Storage (SSD)	1 TB
5ms	Magnetic Disk	8 TB



- Disk is about *ten billion* times larger than registers, but has about *ten million* times larger delay (latency)
- Goal is to work as much as possible in the top levels
- Large, rarely-needed data is stored at the bottom level

Sorting and Binary Search



- We know it's easy to find data if it's in a *sorted* list
 - That's why printed dictionaries and phone books are alphabetical
- **Binary Search** is how computers find entries in a sorted list
 - Let's say you're looking for the word "key" in a list of 10,000 words
 1. Compare "key" to the word in the middle position (5,000th word)
 2. If you're lucky and that middle word is *equal* to "key", then you're done!
 3. If the middle word is *greater than* "key" then go back to step 1, but use just the left half of the list (words 0 through 4,999)
 4. If the middle word is *less than* "key" then go back to step 1, but use just the right half of the list (words 5,001 through 10,000)
 - It will take at most $\log_2 N$ steps to find the entry, where N is the list size
 - e.g., 32 steps for binary search in a list of 4 billion (because $2^{32} \cong 4$ billion)

3

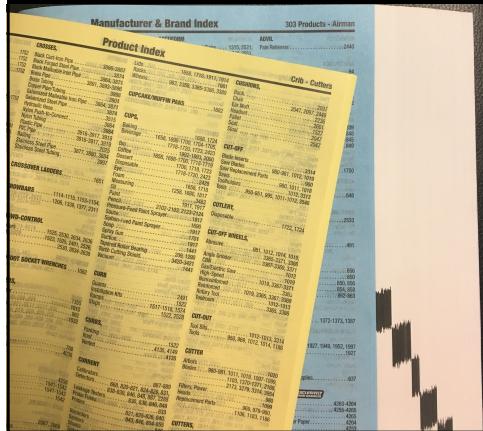
Why sorting is not enough

1. You can't *insert new data* without shuffling everything over to make room
2. You can't sort in *multiple dimensions*
 - Let's say you want to find a product quickly according to either its name, manufacturer, or price. You can only sort by one of the three columns
3. It doesn't take advantage of the hardware's storage hierarchy
 - The binary search will have to access the disk in every step because the index is distributed over the full data set
 - It would be better to put all the index data close together

4

A printed catalog handles multiple indexes by adding more pages

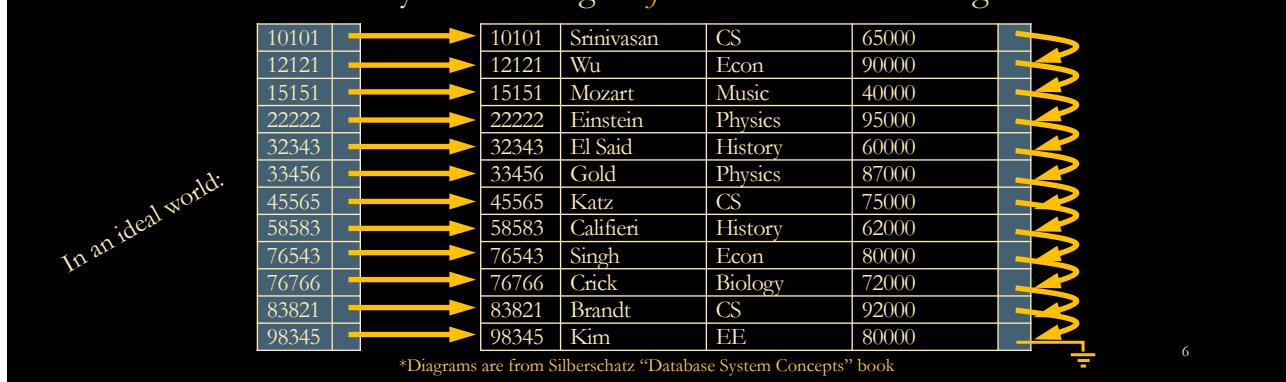
- Real-life example: Grainger catalog
- White pages: data (unsorted)
- Additional yellow pages
 - Index of products
 - Sorted by product type name
- Additional blue pages
 - Index of manufacturers
 - Sorted by manufacturer name
- Easy to locate products and manufacturers
 - ...without sorting the data in either dimension



5

Index data use the storage hierarchy more effectively

- All of the index data may fit in a higher level of the storage hierarchy
- For example, might be able to fit this one column entirely in RAM and avoid using magnetic disk (ns instead of ms access latencies)
- Arrows mean that you're storing a *reference* to another storage location



6

Index data use the storage hierarchy more effectively

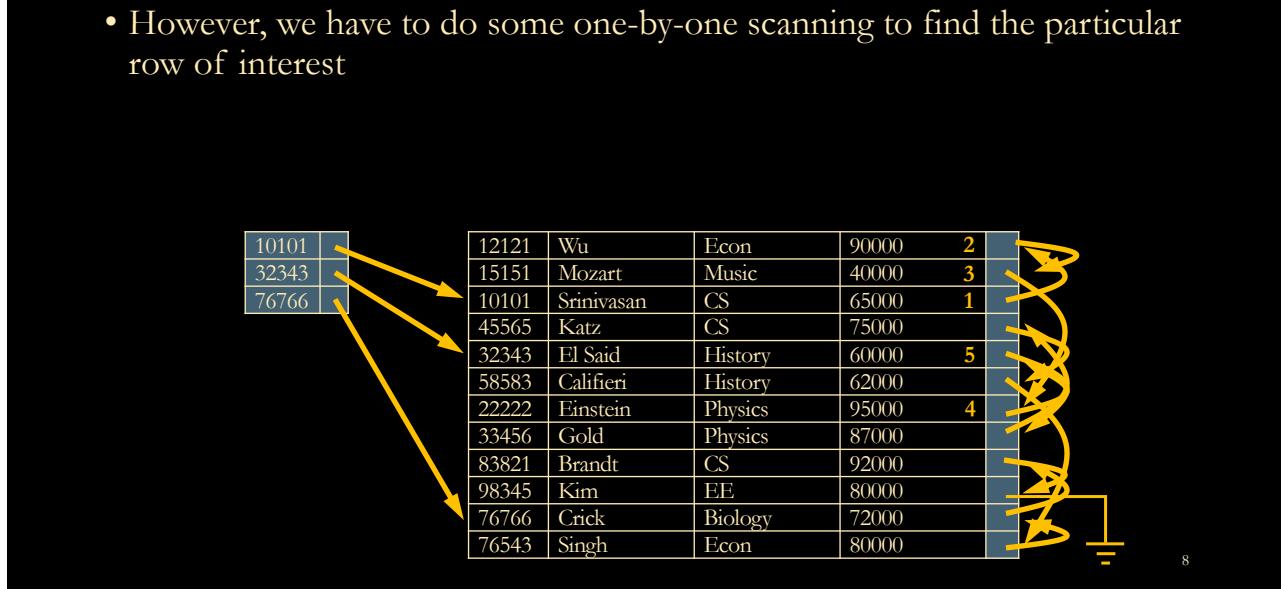
- All of the index data may fit in a higher level of the storage hierarchy
- For example, might be able to fit this one column entirely in RAM and avoid using magnetic disk (ns instead of ms access latencies)
- Arrows mean that you're storing a *reference* to another storage location



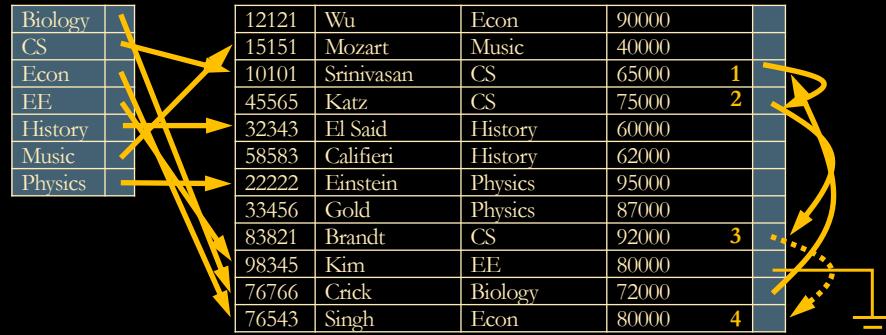
7

A *Sparse* Index may fit in an even higher storage level

- However, we have to do some one-by-one scanning to find the particular row of interest



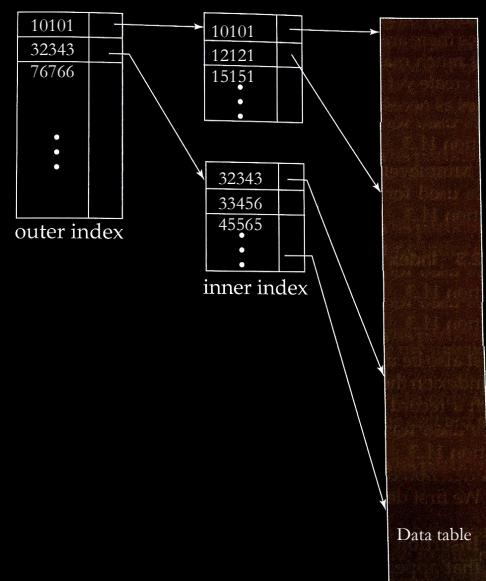
Sparse Index can be used for repeated values



9

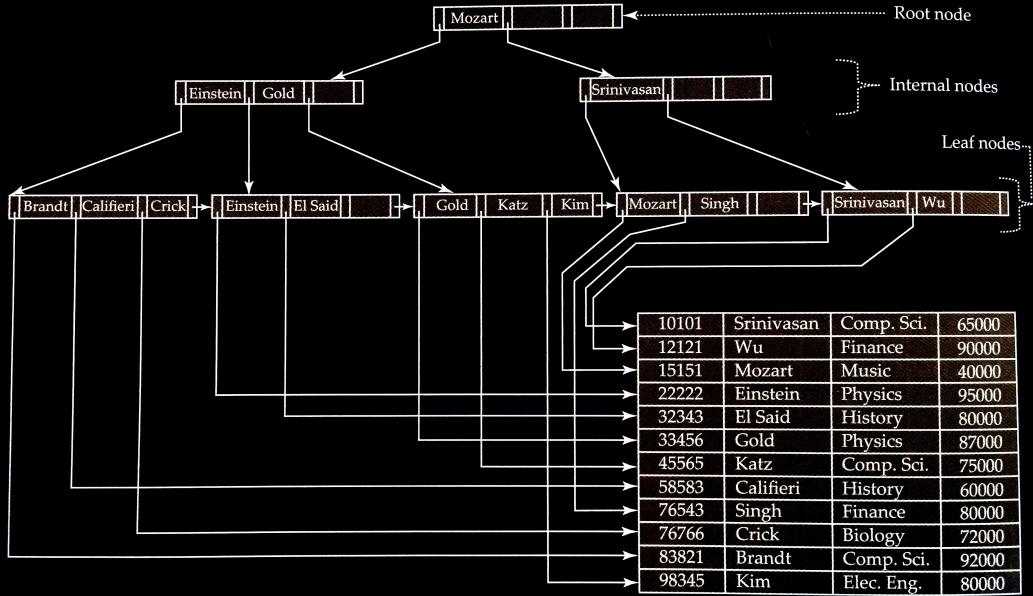
Two-level Index gives us the best of both

- The outer level is small and fits in fast memory, like the sparse index
- The inner level is larger but still fits in a faster memory level than main data, and it tells us exactly where to find data. No need to scan one-by-one



10

B+Trees generalize the concept of multi-level indexes



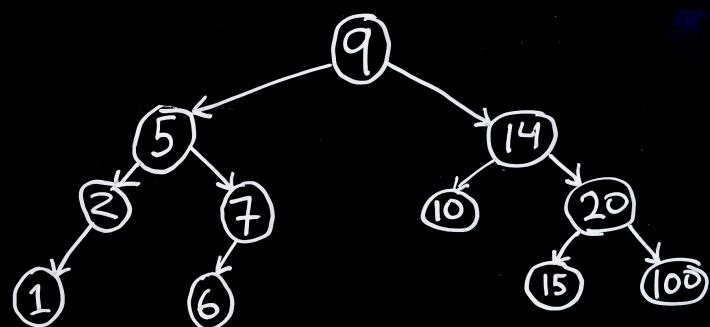
11

Detour: Binary Search Tree

- A “Data Structure” that allows you to search for items quickly.
- Instead of sorting data like this:

1	2	5	6	7	9	10	14	15	20	100
---	---	---	---	---	---	----	----	----	----	-----

- Organize data into a tree like this:



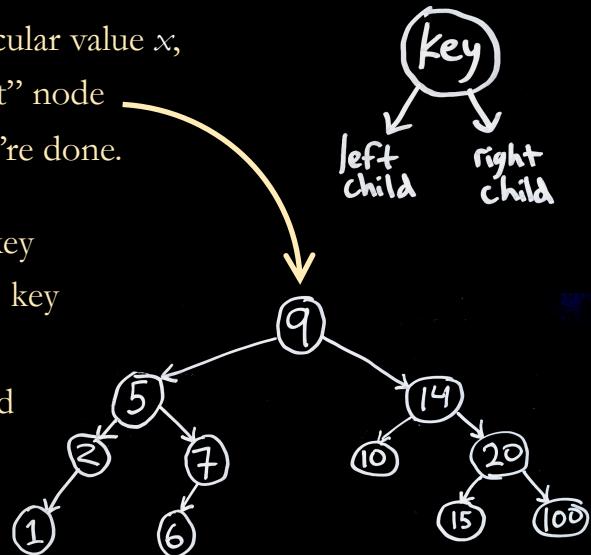
12

Binary Search Tree mechanics

Given that you are searching for a particular value x ,

1. Always start your search at the “root” node
2. If this node’s key equals x , then you’re done.
You have found it!
3. Move to the *left child** if $x <$ node’s key
4. Move to the *right child** if $x >$ node’s key
5. Go to step 2

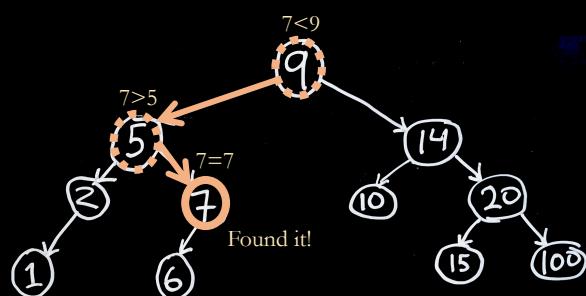
* In steps 3 & 4, if there is no such child
then the key is absent



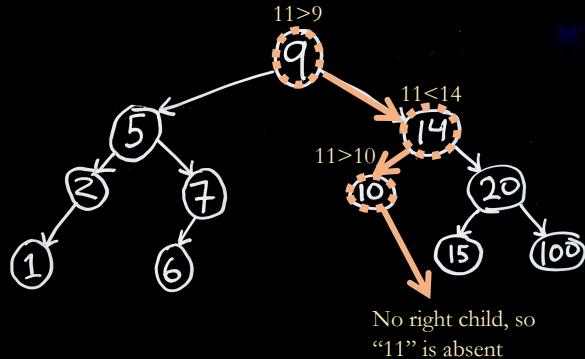
13

BST Search Examples

Searching for “7”:



Searching for “11”



14

Building a Binary Search Tree

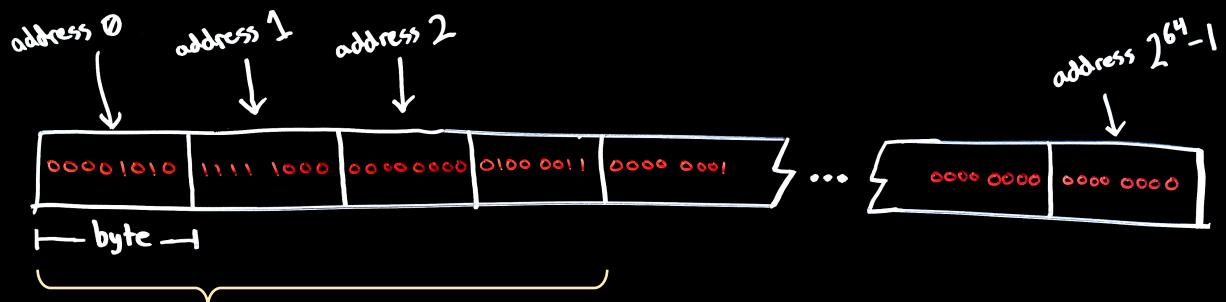
- Works just like search
- Find where key *should be* then add it

So, let's talk about memory...

15

Computer memory is basically just a big array

- All data is stored as a sequence of zeros and ones
- Central Processing Unit (CPU) specifies which *address* to read and write
 - Addresses refer to a byte-sized sequence of bits (8 bits)



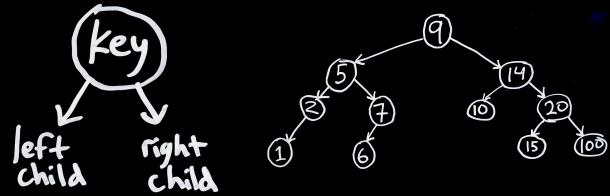
These 4 memory addresses could represent many different things.
For example, a single 32-bit float, or a few characters of UTF-8 text.

16

A Tree is a type of “graph”
A collection of nodes referring to other nodes



Address	Key	Left Child	Right Child
101	20	310	105
102	5	104	103
103	7	108	NULL
104	2	107	NULL
105	100	NULL	NULL
106	9	102	891
107	1	NULL	NULL
108	6	NULL	NULL
...			
310	15	NULL	NULL
311	10	NULL	NULL
...			
891	14	311	101

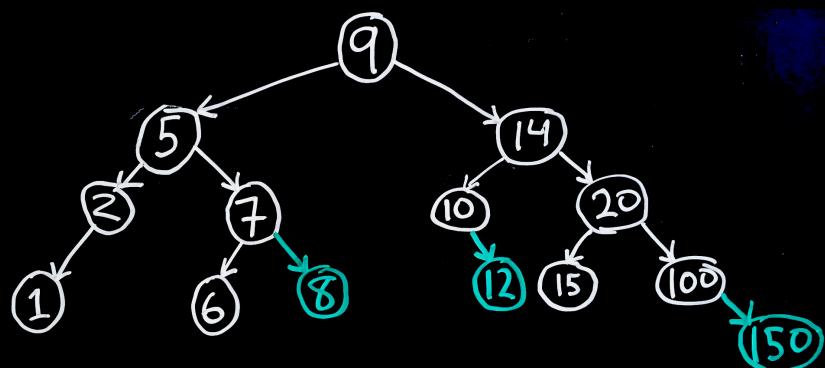


- In this case the root is at address 106
- It's a *linked* data structure
 - Allows us to use a one-dimensional array of computer memory to represent something with complex geometry
 - The tree can be reshaped very quickly by just changing a couple of child values

17

Quick insertions and deletions in a tree

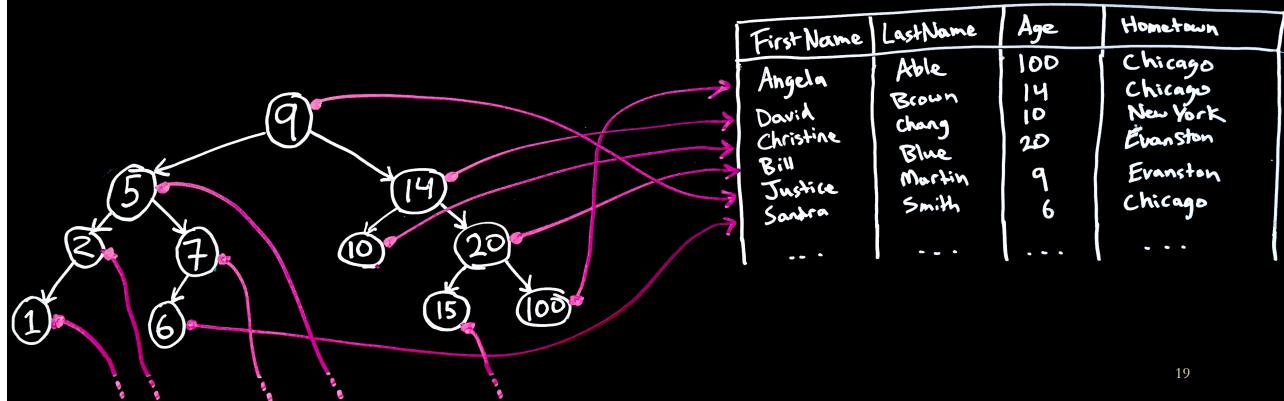
- There is no need to shift all data, just change a few nodes



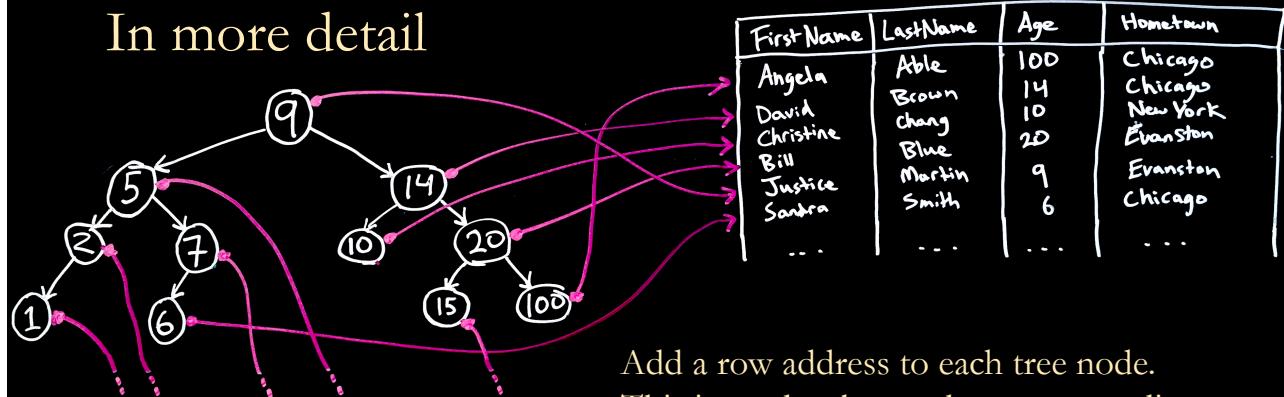
18

Using a BST as an index

- Add one piece of additional info to each node – a row address
- Below, a BST is an index on a table's *Age* column
- Search for an age in the tree, then the tree node directs you to a row



In more detail

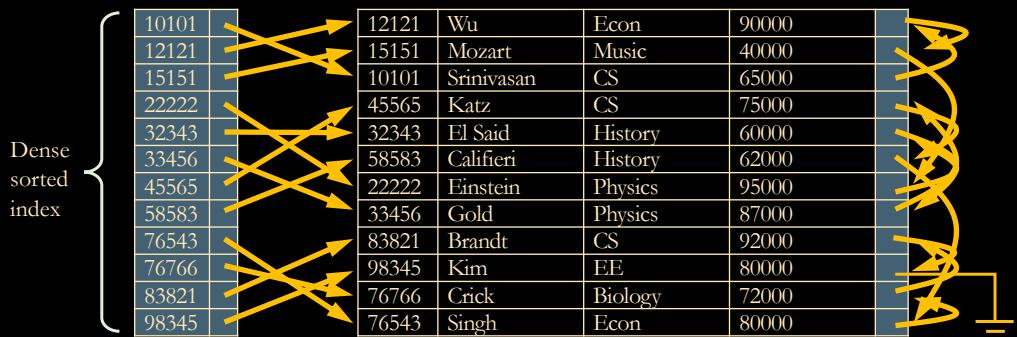


Tree representation in memory:

Address	Key	Left Child	Right Child	Row Address
101	20	310	105	10500
102	5	104	103	10200
103	7	108	NULL	10300
104	2	107	NULL	20200
105	100	NULL	NULL	20800

Reminder: Dense sorted index

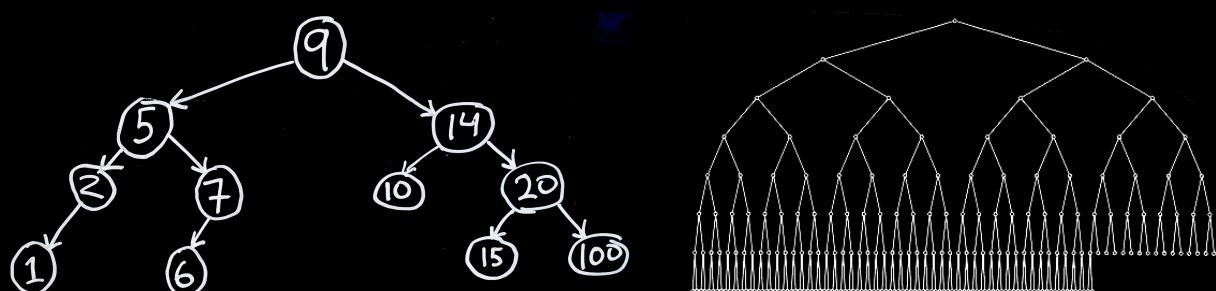
- Cannot add or delete items without shifting data
- Might think about leaving some open space, but it's impossible to predict exactly where space will be needed, and this will change over time



21

Searching a BST can be fast if you eliminate half of the remaining options in each step

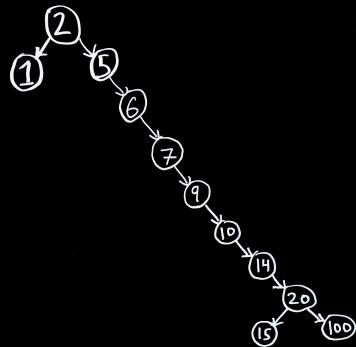
- Like Binary Search of a sorted list, \propto steps lets you scan 2^x data values
- Each additional level can index twice as much data



22

But an *unbalanced* tree is very inefficient

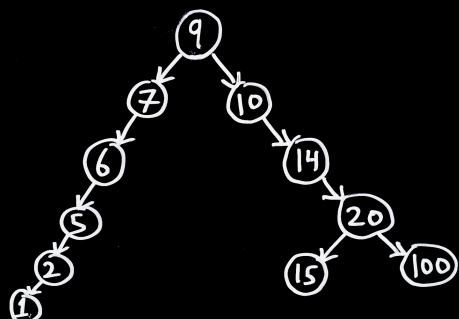
- In the worst case, you would look at every node to find what you're looking for
- The problem is that each step is not eliminating much candidate data
 - The “path not taken” is empty



23

Maintaining balance in a tree

- It's not enough to just choose your root node wisely. This tree is better, but still not well balanced:

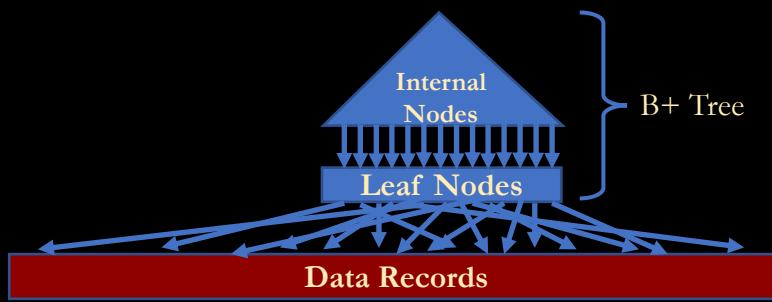


- *Every node* must have a value near the middle of all its descendants

24

Computer scientists have invented many different self-balancing trees:

- For example: red-black trees, AVL trees, splay trees, and b-trees
- They all require some quick maintenance work to restore balance after each insertion or deletion.
- Databases mostly use B⁺-trees for indexes
 - We won't go too deep into the details of B⁺-trees in this course



25

